

Building ETL Pipelines from Streaming Data with Kafka and KsqlDB

UNDERSTANDING KSQLDB AND KAFKA STREAMS



Eugene Meidinger


BI CONSULTANT

@sqlgene

www.sqlgene.com



Overview




Topics



Deployment



Aggregations



Enrichment



Challenges



Kafka wasn't designed
for stream processing.



Kafka is an event streaming platform

Apache Software Foundation



Kafka is an event streaming platform



Kafka is an event streaming platform



Kafka is an event streaming platform



Benefits of Kafka

Decoupling



Scaling



Space



Time


What Is Kafka Streams?



Kafka Streams is a library
that lets you speak
in streams, not events.



What Does Kafka Streams Provide?



Nouns



Verbs



Scalability



Reliability



Aggregating Data


Kafka Topic



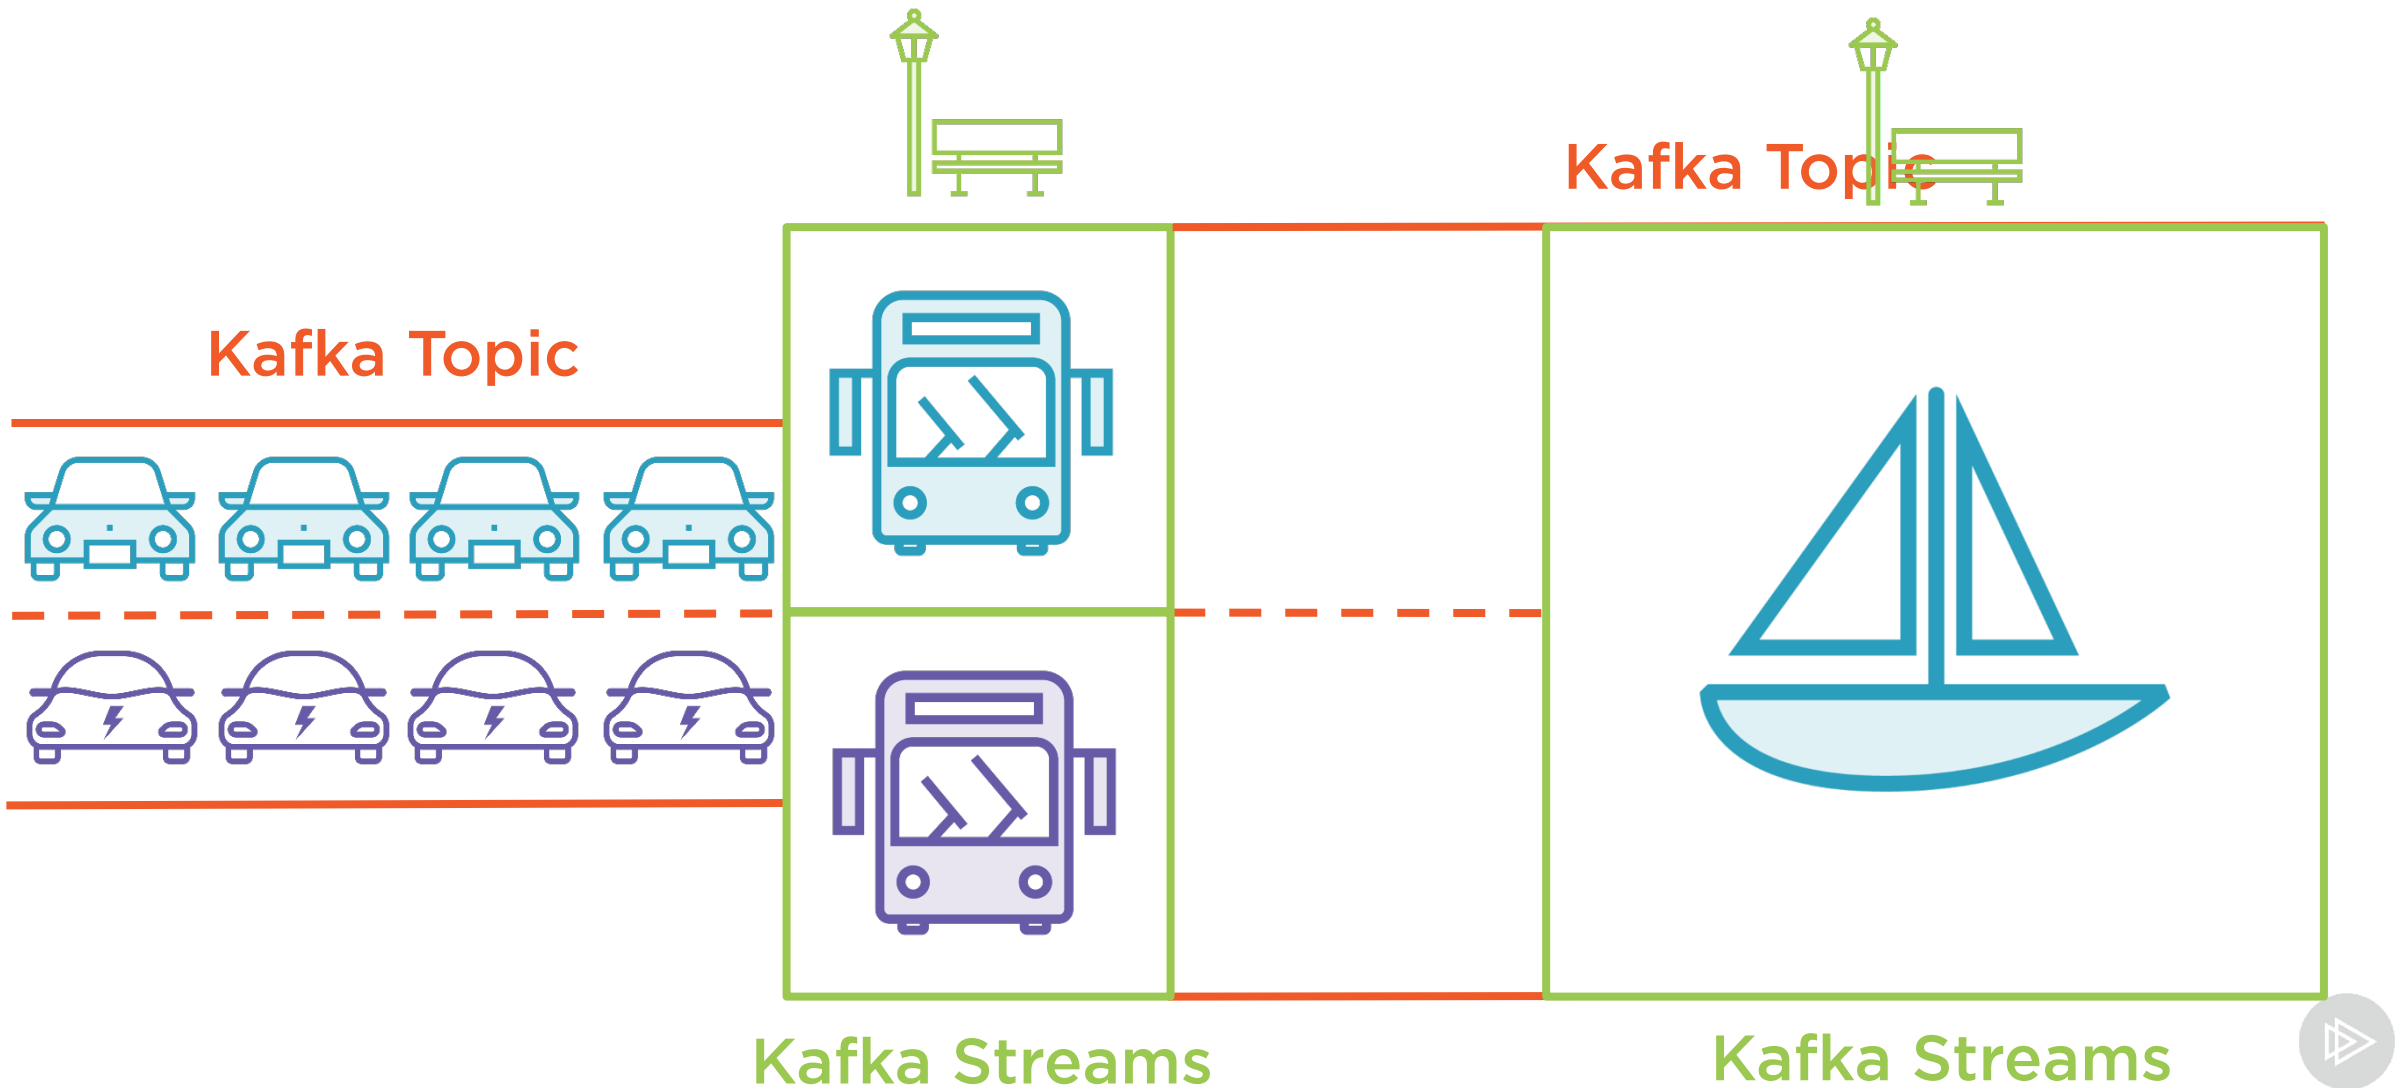
Aggregating Data



Application Instances




Chaining Processors



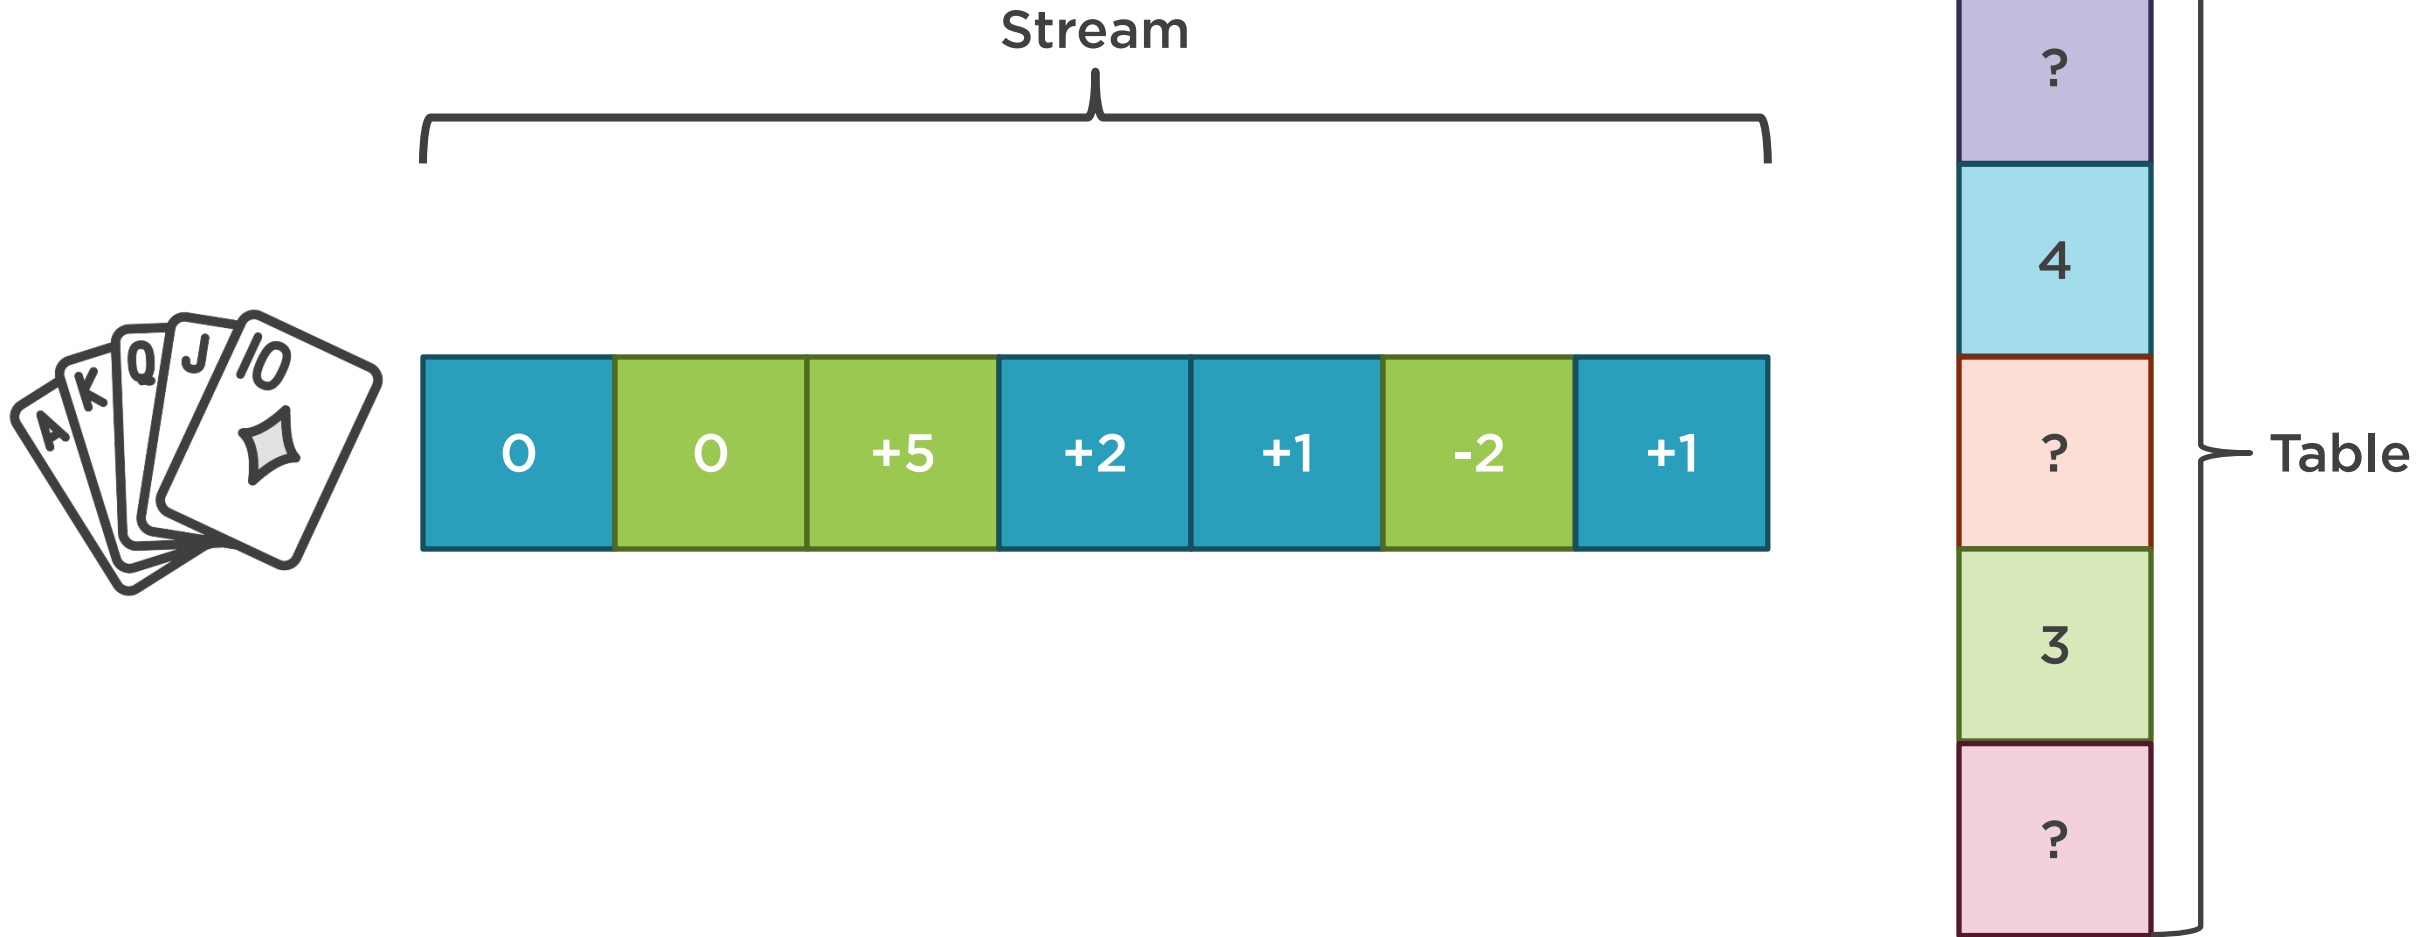
Processor Topology



Streams and Tables



Streams and Tables



What is KsqlDB?



“ksqlDB is an event streaming database purpose-built to help developers create stream processing applications on top of Apache Kafka.”

Confluent



“ksqlDB is an event streaming database [...] on top of Apache Kafka.”

Confluent



What is MYSQL?


What is Postgres?



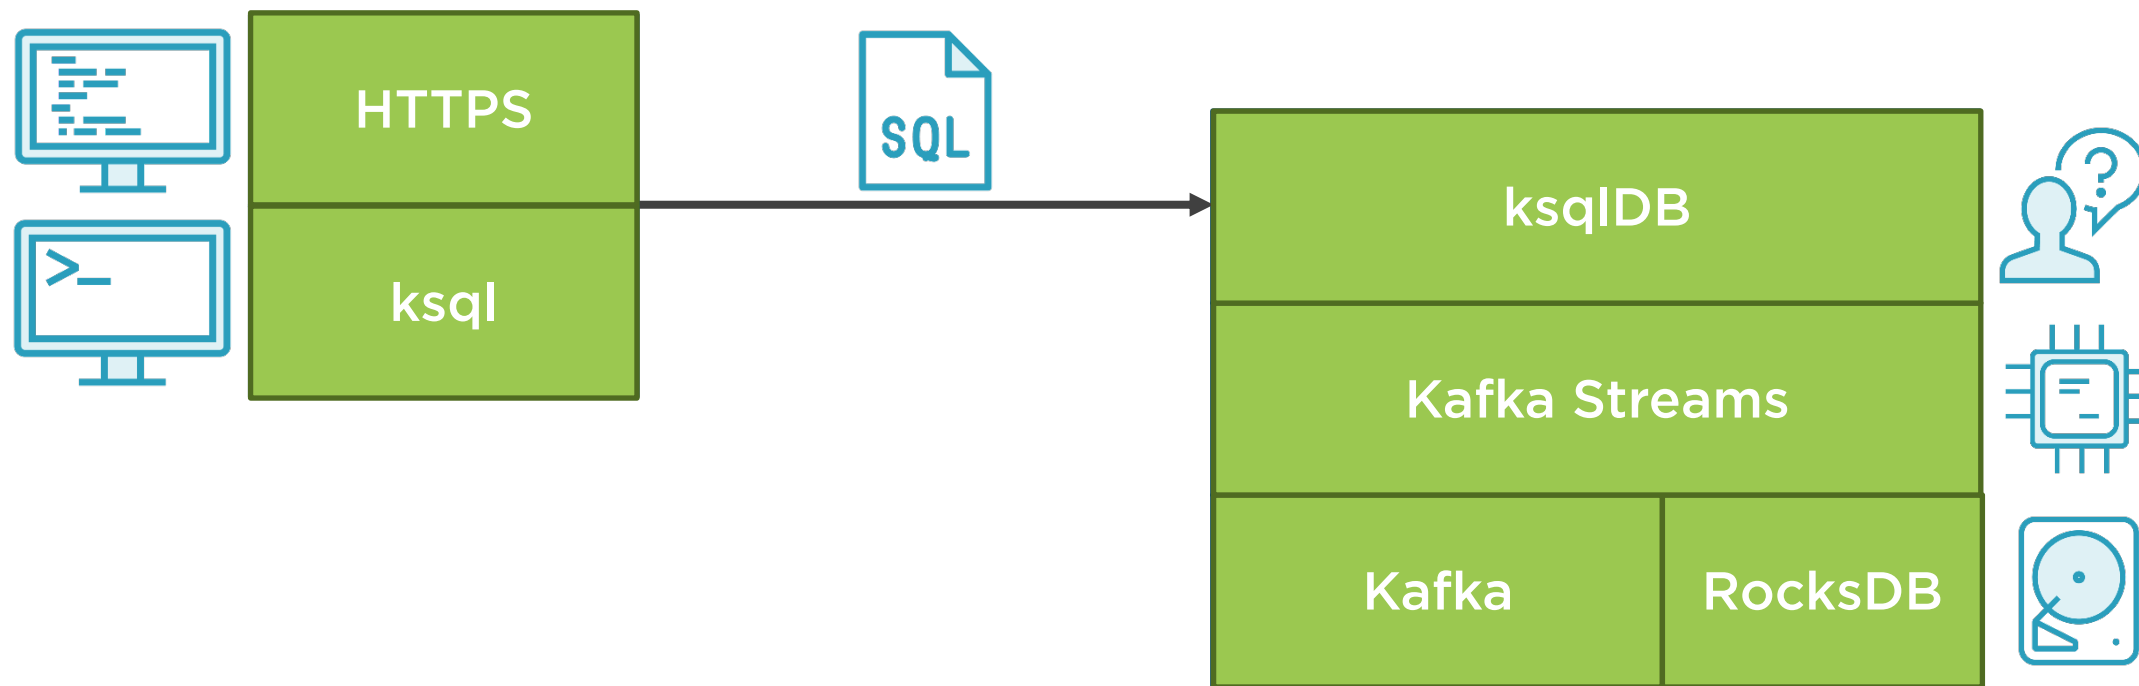
What Is a Database?



What Is a Database?




What Is a Database?



What Is KsqlDB?




SQL Variant



Command Line
Interface



Web Server



Top Layer



Overview



Kafka stores events

- Events are small historical records

Kakfa Streams simplifies ETL

- Uses streams and tables
- Manages scaling out

ksqlDB gives us an SQL interface

- No compiling, less work

