

Writing and Deploying KsqlDB Queries



Eugene Meidinger

BI CONSULTANT

@sqlgene www.sqlgene.com



Overview



SQL + streaming features

Creating a query

Deployment options

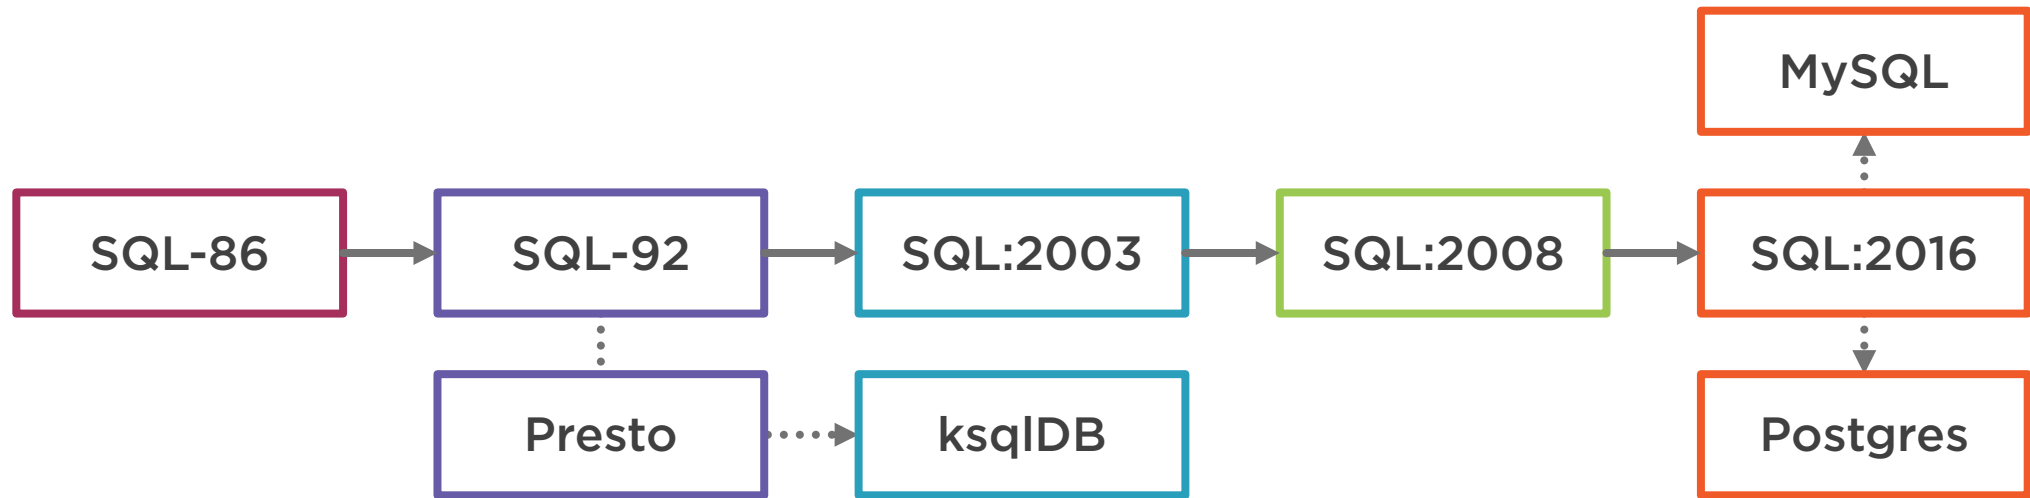
- CLI
- Web server
- Headless deployment




KsqlDB queries are ANSI SQL plus extra features.




SQL Family Tree



Key Differences



Table




Stream



Modify



Append



Pull Query




Push Query

Creating Your First KsqlDB Query



Weather Analysis



Hundreds of weather sensors

Thousands of readings

Stored in Kafka

Analysis with ksqlDB

Types of Queries



Push Query



Pull Query



Joins

Push Query



Pull Query



Join Query



KsqlDB let's us modify
and summarize data.



Interactive versus Headless Deployments



Headless just means
no front-end.



Deployment Modes



Kafka Streams



Headless



Interactive



Comparison Recap

Headless

File-based

Predictable workloads

Production environment

Interactive

Query-based

User-defined workloads

Development environment



Demo



Running the ksqlDB CLI

Creating a push query

Creating a pull query

Joining the results



Deploying a Streaming ETL Pipeline in Headless Mode




Streaming data terms
are confusing.



Streaming Applications




Streaming ETL Pipeline




Event-driven Microservice



Streaming Applications




Transforms Data



Event-driven Microservice



Streaming Applications



Transforms Data



Triggers Side-effects



Deploying in Headless Mode



Kafka Connect



Queries File



Command Line Parameter (--queries-file)



Kafka Client



Additional Resources



Kafka Connect Fundamentals
By Bogdan Sucaciu



Streaming ETL tutorial
<https://docs.ksqldb.io/en/latest/tutorials/etl/>



Summary



Flavor of SQL

Types of queries

- Push query
- Pull query

Types of deployments

- Interactive
- Headless

