Data Wrangling in Python for Machine Learning Engineers

GETTING STARTED

Mike West
MACHINE LEARNING ENGINEER

Module Overview

What is supervised learning?

What is the machine learning process?

Why clean data is so important?

Skills for this course

Wrangle your data set in Python

Overview

The Pandas Dataframe

Titanic Dataset

Model our wrangled data set

Summary

Your Skills

Not Required

Machine Learning

Deep Python knowledge

Statistics or advanced math

Required

Basic Python

Tables, columns and rows

Basic math and statistics

Why Take this Course?

Required real-world skill

In the real-world, machine learning engineers spend most of their time wrangling data.

Why Take this Course?

- Required real-world skill
- Applied machine learning is data wrangling
- Models need well cleansed numerical data
- Improved model performance

Two Types of Machine Learning

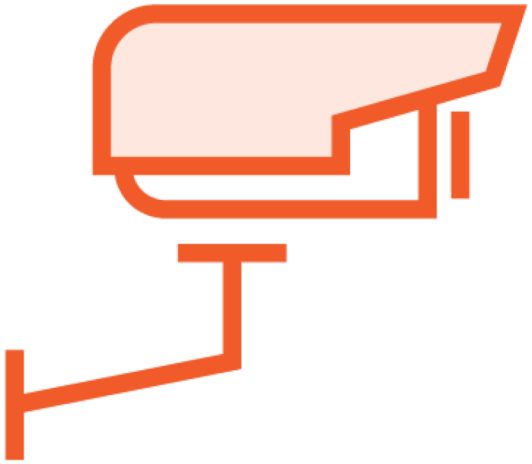

Supervised

The models are fed clean labeled numerical data.

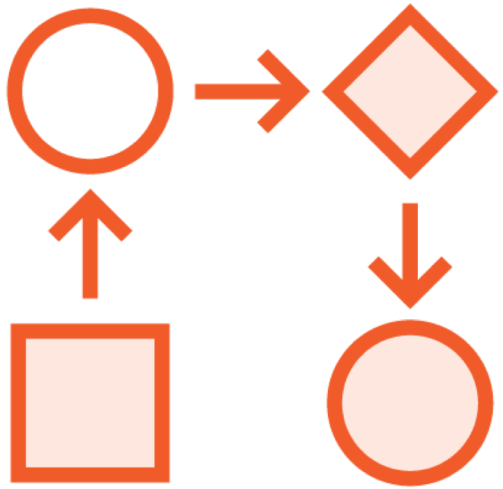

Unsupervised

The models find patterns and structure in unlabeled data.

Supervised Machine Learning

	0	1	2	3	4
0	1	Braund, Owen	male	52	0
1	2	Cumings, John	male	38	0
2	3	Heikkinen, Laina	female	26	1
3	4	Futrelle, Jacques	male	35	1
4	5	Allen, Henry	male	35	0

	0	1	2	3
0	1	Braund, Owen	male	52
1	2	Cumings, John	male	38
2	3	Heikkinen, Laina	female	26
3	4	Futrelle, Jacques	male	35
4	5	Allen, Henry	male	35

Unsupervised Machine Learning

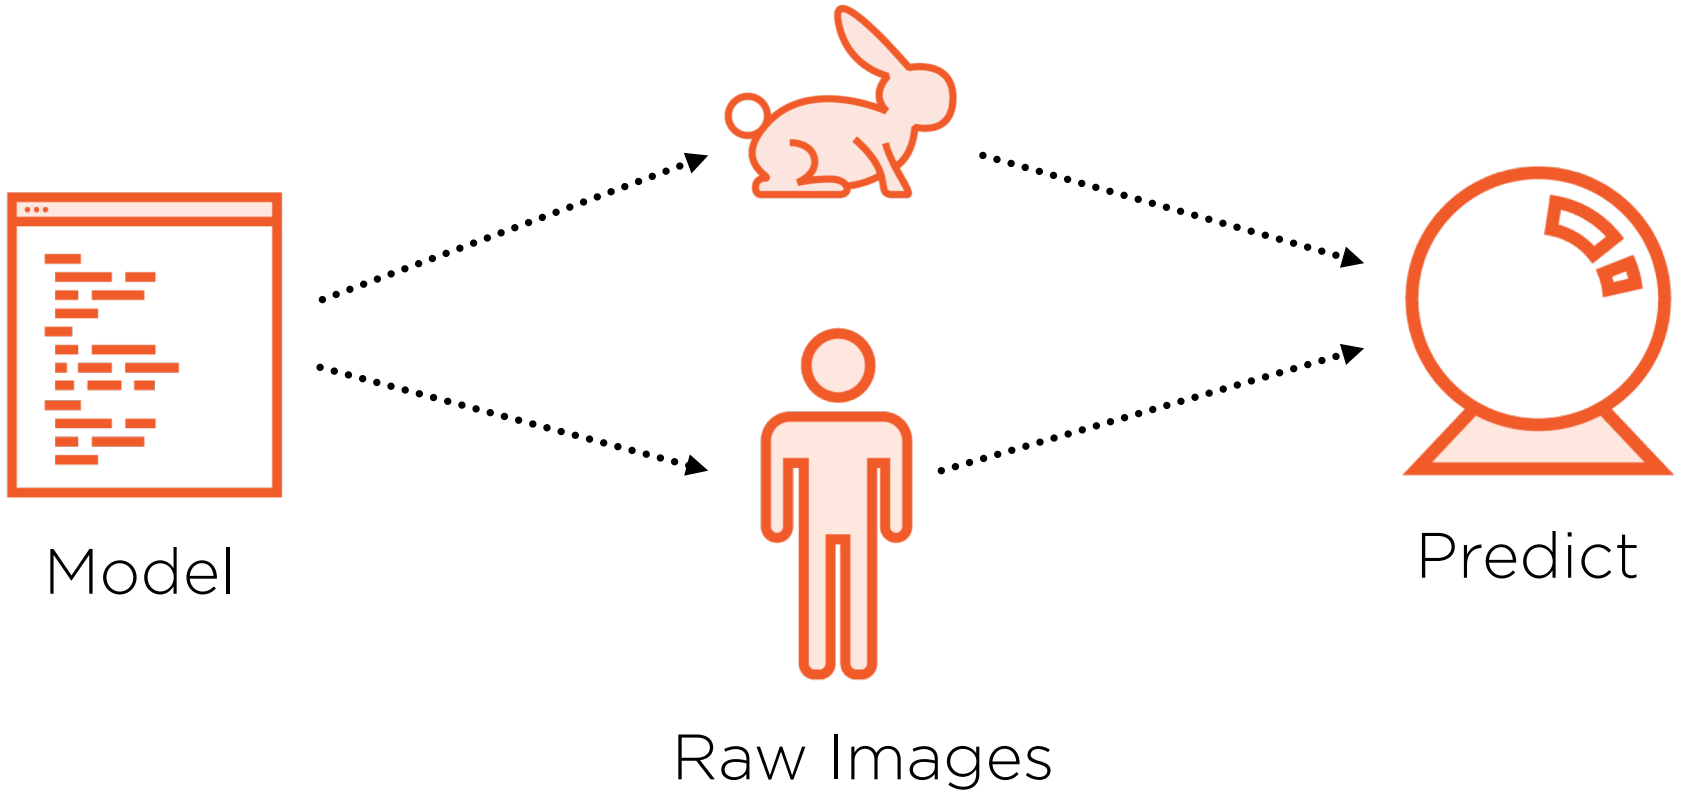

The Machine Learning Process

Most data is sourced from relational databases. Data is also taken from big data platforms and comma separated files.

Data wrangling is time consuming and difficult. In Python, Pandas is the preferred tool.

Modeling involves feeding data to traditional algorithms and artificial neural networks.

Once the model has been tuned and tested it’s ready for production. The models should perform well on data they’ve never seen.

Sourcing Our Data

Databases
Most models are currently sourced from relational data

Flat Files
Most open source data are comma separated files

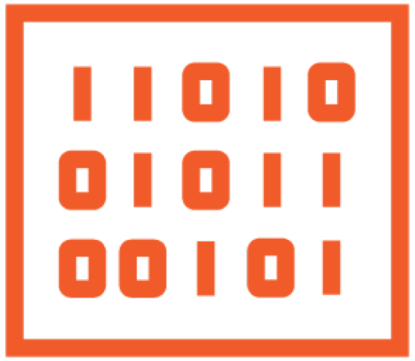

Big Data
Approximately 90% of data collected globally is unstructured

The Core Algorithms

Traditional Models

All models that are not artificial neural networks

Neural Networks

Models that are artificial neural networks

Final Phase of Machine Learning

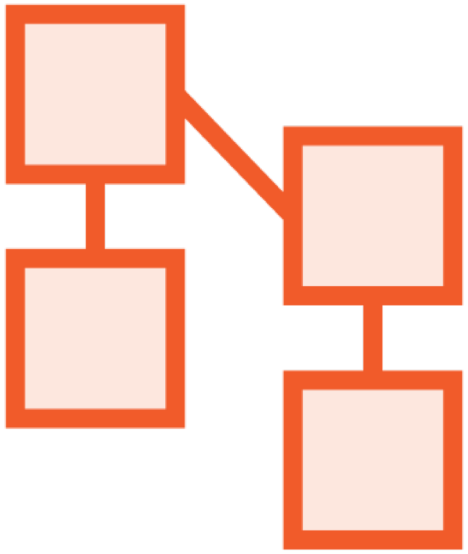

Completed Model

The model has been trained on our dataset. We've also tuned the model for best performance.

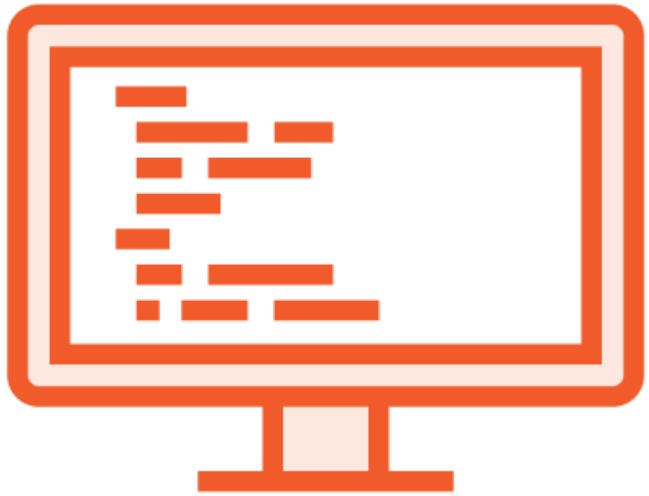

Production

The model is live. The model should be able to make accurate predictions on data it has never seen.

Demo

Python 3.6

Jupyter Notebooks

Import Pandas library

Remove an attribute

View wrangled dataset