

Reducing Complexity in Nonlinear Data



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Manifold learning for dimensionality reduction

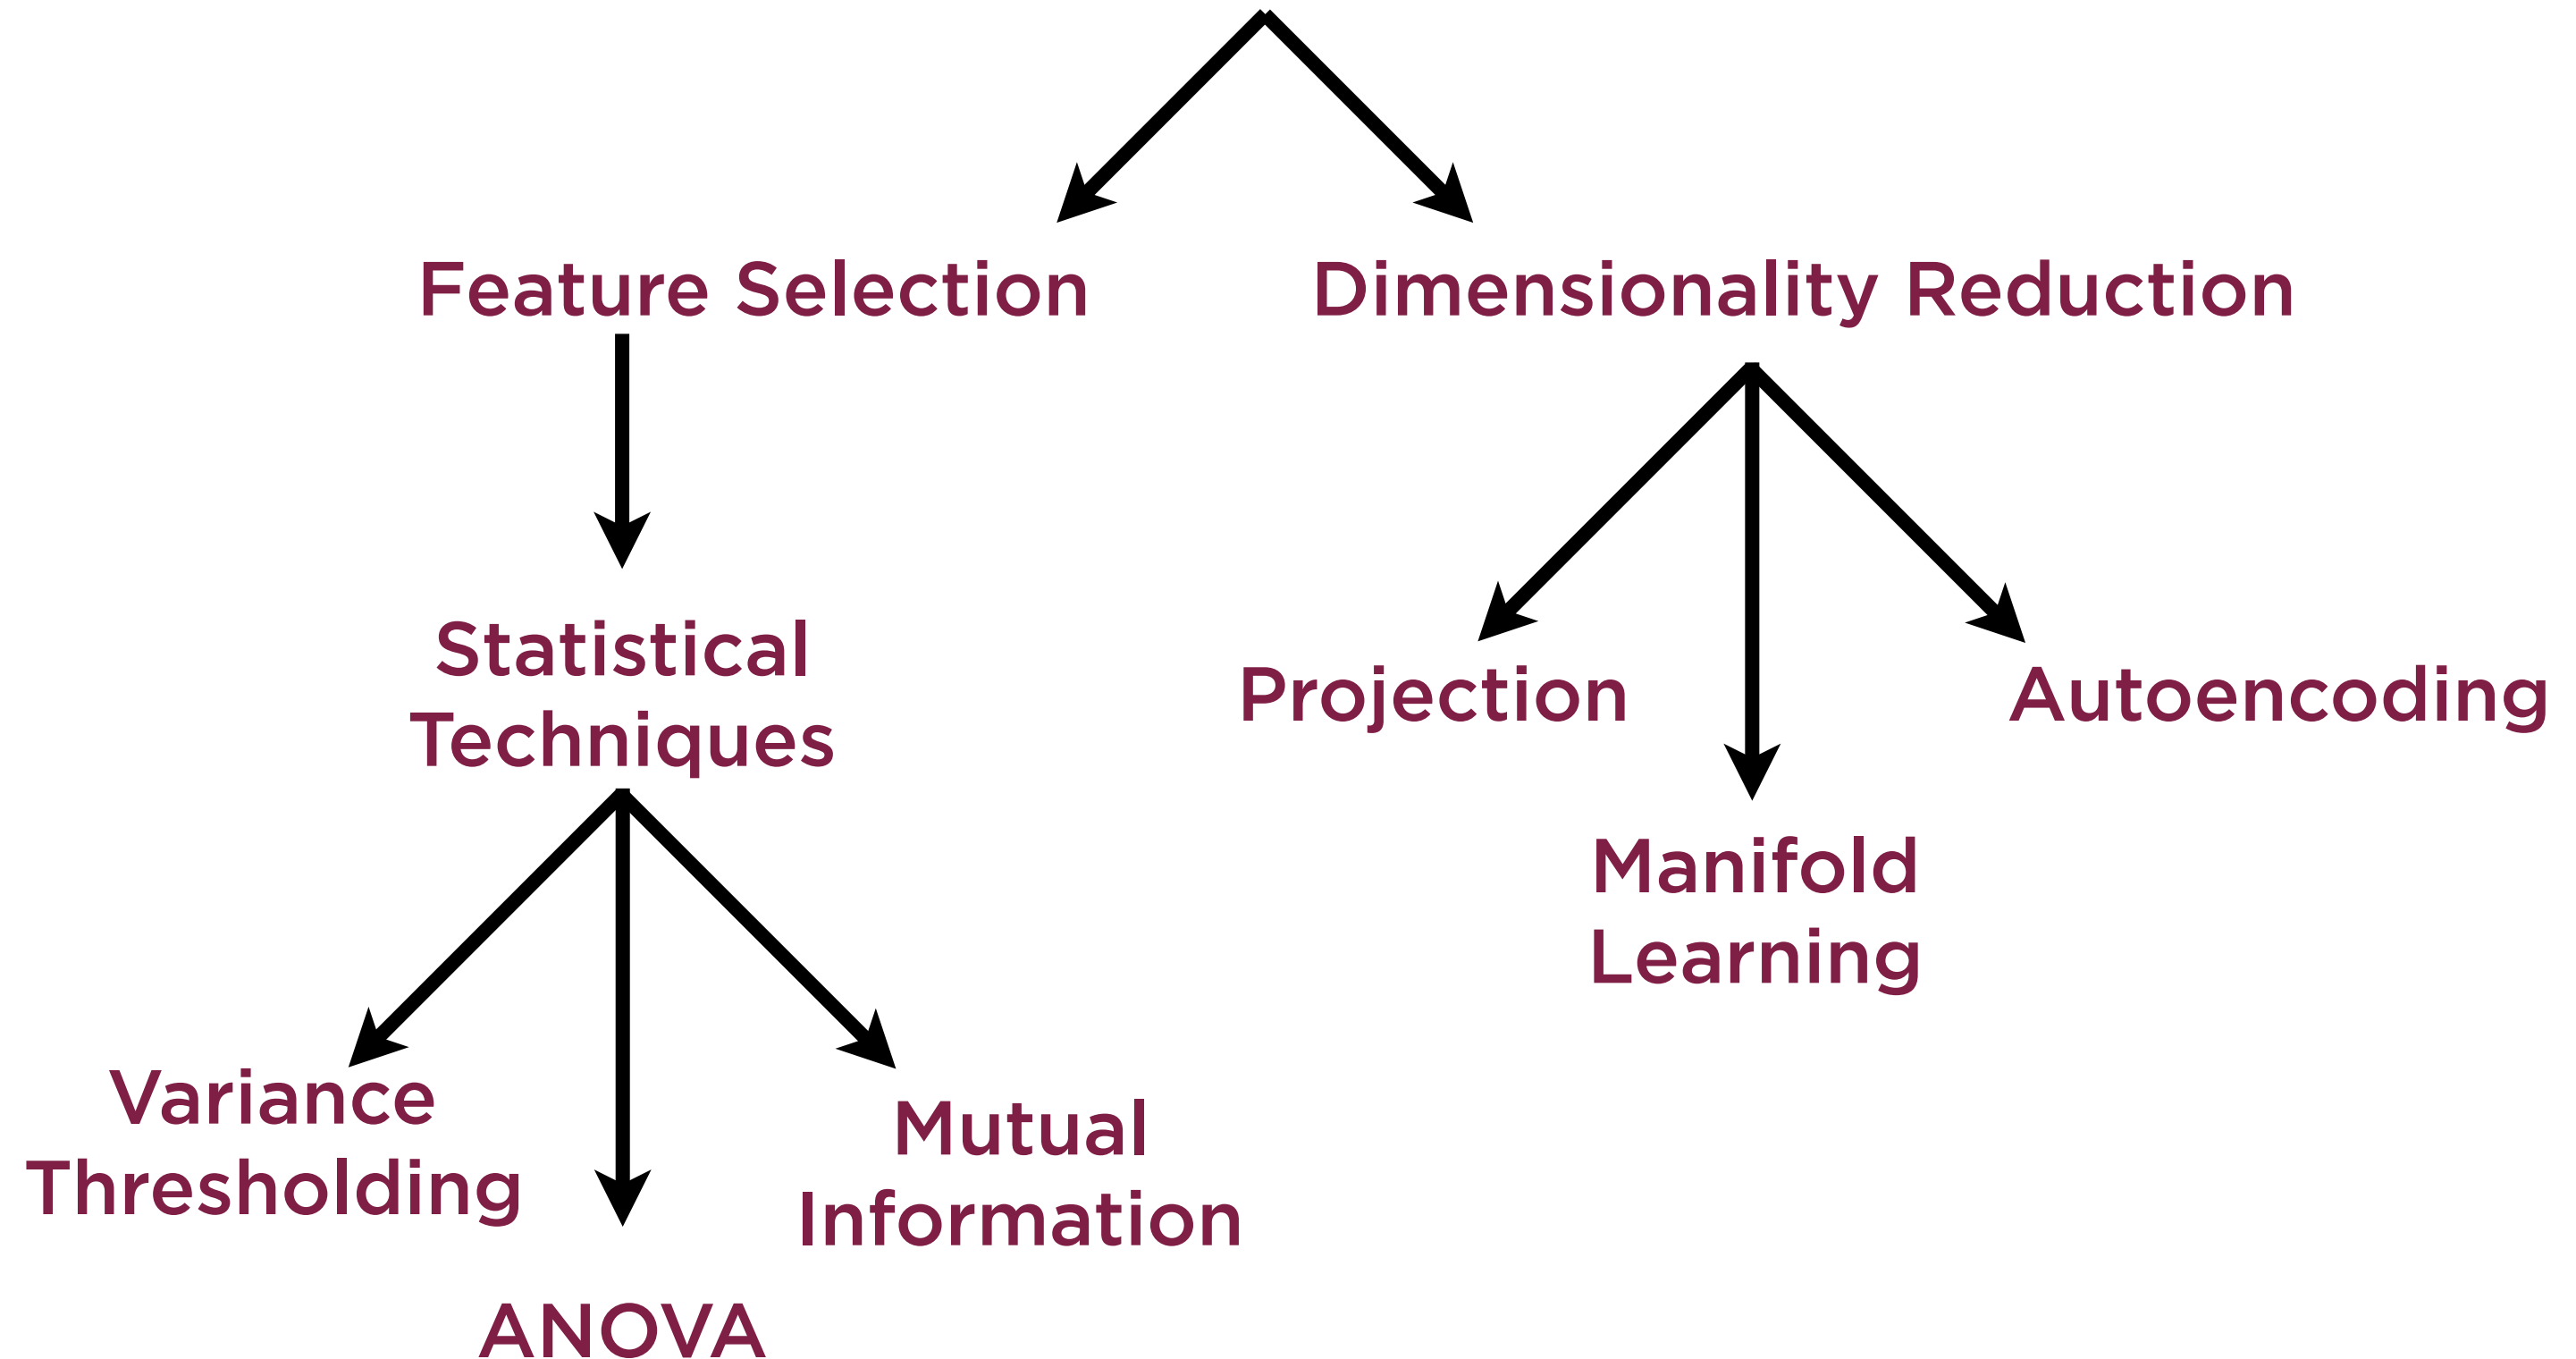
Implementations of manifold learning

Multidimensional scaling, Isomap, Locally Linear Embedding


t-Distributed Stochastic Neighbor Embedding (t-SNE)

Kernel PCA

Reducing Complexity




Reducing Complexity



Unroll the data so that twists
and turns are smoothened out



Reducing Complexity




Works best when data lies along a
rolled-up surface such as a Swiss
Roll or S-curve

**Manifold
Learning**



Reducing Complexity



e.g. MDS, Isomap,
LLE, Kernel PCA

**Manifold
Learning**

Manifold Learning

Manifold Hypothesis:
Very complicated data is often not
that complicated after all

Choosing Manifold Learning

Use Case

Y not linearly related to X

Very high dimensionality of X (e.g.
pixel counts in image data)


Many constraints on allowable
values of X-variables (sparse
features)

Three-dimensional plots of Y
against pairs of X indicate
manifold shape


Possible Solution

Manifold learning


Linear Data




Nonlinear Data



Manifold Data




Manifold Hypothesis



Many high-dimensional datasets can be easily unrolled so that they lie along a much lower dimensional manifold


Manifold Hypothesis



High-dimensional
data



Manifold Learning



Low-dimensional
embedding

Manifold Learning Techniques

MDS

Isomap

Kernel PCA

LLE

t-SNE

Spectral Embedding

Multidimensional Scaling (MDS)

Aims to preserve pair-wise Euclidean distances between all points while reducing dimensionality. Some intuitive similarities to MSE regression in underlying math.

Isomap

Aims to preserve pair-wise Euclidean distances between neighboring points only (not all points) while reducing dimensionality; works out equivalent to preserving geodesic distance between all points.

Locally Linear Embedding

Expresses each point as centroid (weighted average) of nearest neighbors; then tries to maintain same weights upon conversion to new dimensions.

Spectral Embeddings

Builds a graph where each point serves as a node; then fits a smooth function in lower dimensional space to pass through all nodes. Often implemented using technique called Laplacian Eigenmaps.

t-distributed Stochastic Neighbor Embedding (t-SNE)

Aims to keep similar points together and dissimilar points apart. First fits a Student-t probability distribution to the data, hence the name. Widely used in visualizing clusters.

Kernel PCA

First apply technique called the kernel trick to go to even higher (infinite) number of dimensions, then perform PCA to come down to very low-dimensional space.

Demo

**Implement Manifold Learning using a
Swiss Roll dataset**

Demo

Implement Kernel PCA

Summary

Manifold learning for dimensionality reduction

Implementations of manifold learning

Multidimensional scaling, Isomap, Locally Linear Embedding

t-Distributed Stochastic Neighbor Embedding (t-SNE)

Kernel PCA